

# Databricks for the SQL Developer

Gerhard Brueckl

Cloud Data Architect @ [paiqo GmbH](#)



SSAS  
MAESTRO  
by Microsoft



@gbrueckl



blog.gbrueckl.at



[gerhard@gbrueckl.at](mailto:gerhard@gbrueckl.at)



<https://github.com/gbrueckl>



[DatabricksPS](#)



[Databricks VSCode](#)



[PowerBI Connector](#)



[www.paiqo.com](http://www.paiqo.com)



# Agenda

What is Databricks / Spark?

How is Databricks / Spark different to traditional RDBMS?

SQL with Databricks

Delta Lake

Advanced SQL techniques

# What is Databricks / Spark?

# What is Spark ?

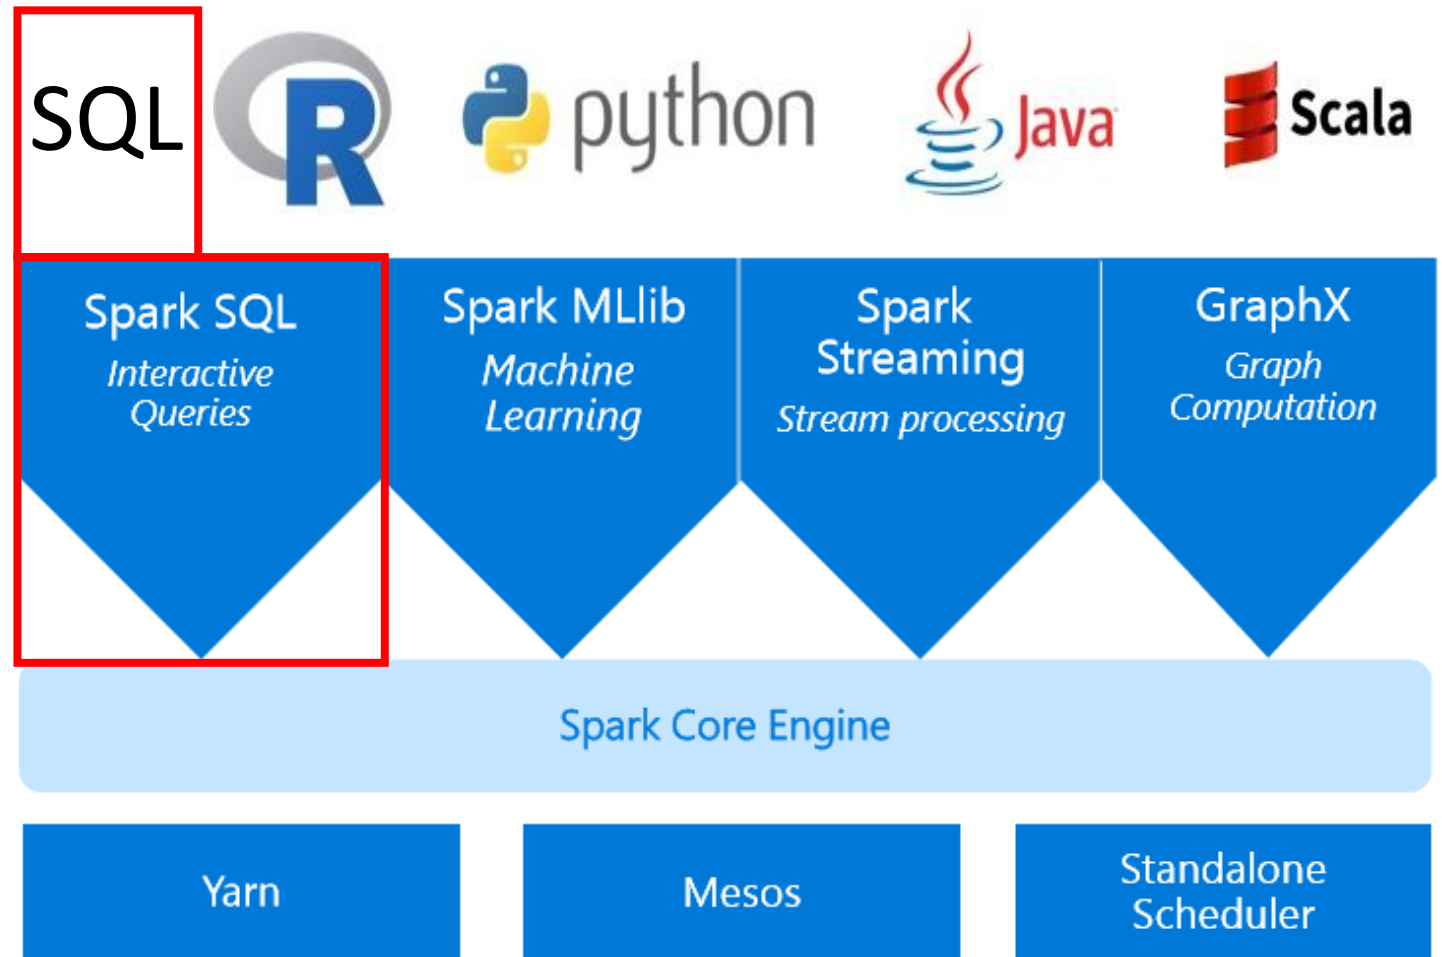
- Open-source Cluster computing framework
  - Massive Parallel Processing with linear scale
- Built for: Speed/Scalability, Ease-of-Use, Extensibility
- Support for multiple languages
  - Java, Scala, Python, R, **SQL**



# What is Spark ?

## Spark unifies

- Batch Processing
- Interactive SQL
- Real-time processing
- Machine Learning
- Deep Learning
- Graph Processing



# What is Databricks ?

- Company that provides a Big Data processing solutions in the Cloud using Apache Spark
  - Databricks on AWS
  - Azure Databricks
  - Databricks on Google Cloud
  - **No on-prem solution!**
- Creators of Apache<sup>®</sup> Spark<sup>™</sup>



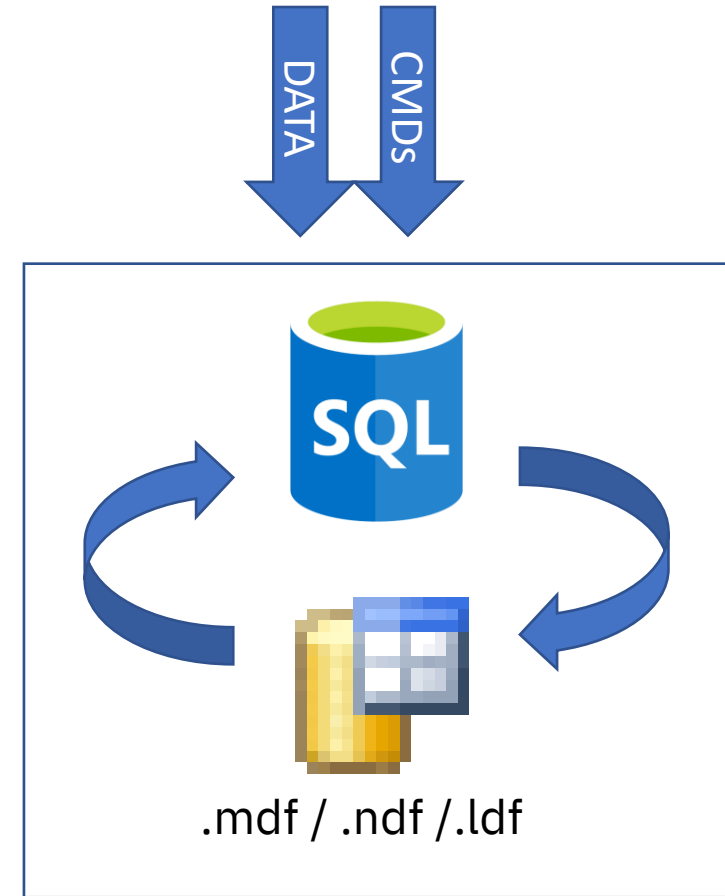
How is Big Data / Spark  
different to traditional RDBMS?



# Classic RDBMS vs. Big Data/Spark

## Classic RDBMS

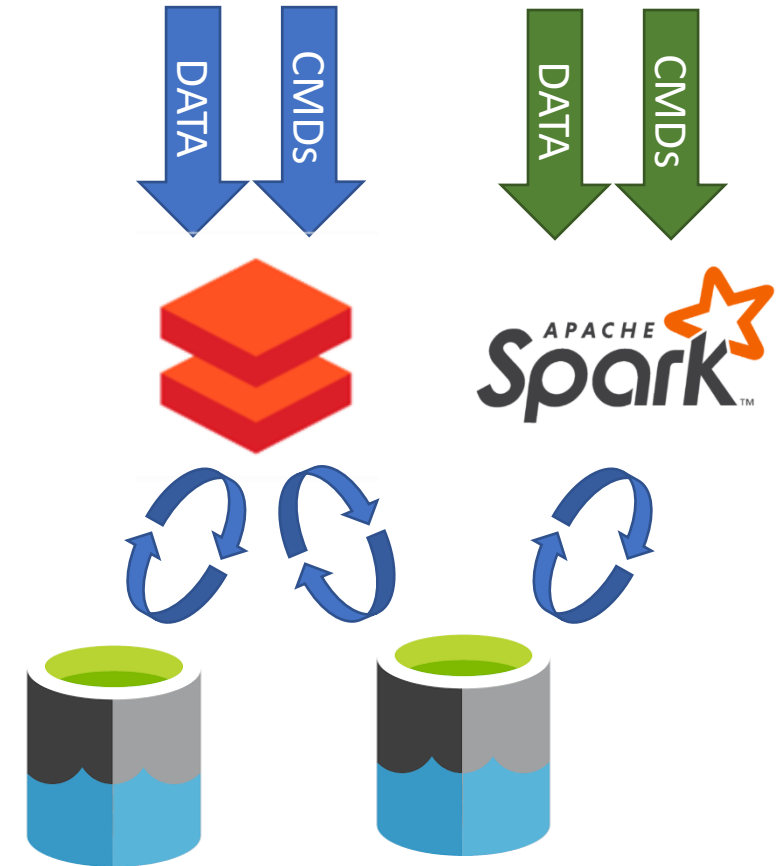
- Single point of access to your data
- Limited resources / no scale-out
- All process use the same resources
  - ETL vs. user queries
- Storage is managed internally



# Classic RDBMS vs. Big Data/Spark

## Big Data processing with Spark

- Separation of storage and compute (!)
- Only spin up compute when necessary
- Can use multiple compute engines
- Cheap storage
- Can attach any storage



# Classic RDBMS vs. Big Data/Spark

RDBMS	Big Data / Spark
Server	Cluster(s) + Metadata/Metastore
Database (e.g. AdventureWorks)	Metastore
Schema (e.g. dbo)	Database
Table (e.g. DimProduct)	Files/Location + Metadata
Index	---
Stored Procedure	Notebook
UDF	UDF
View	View

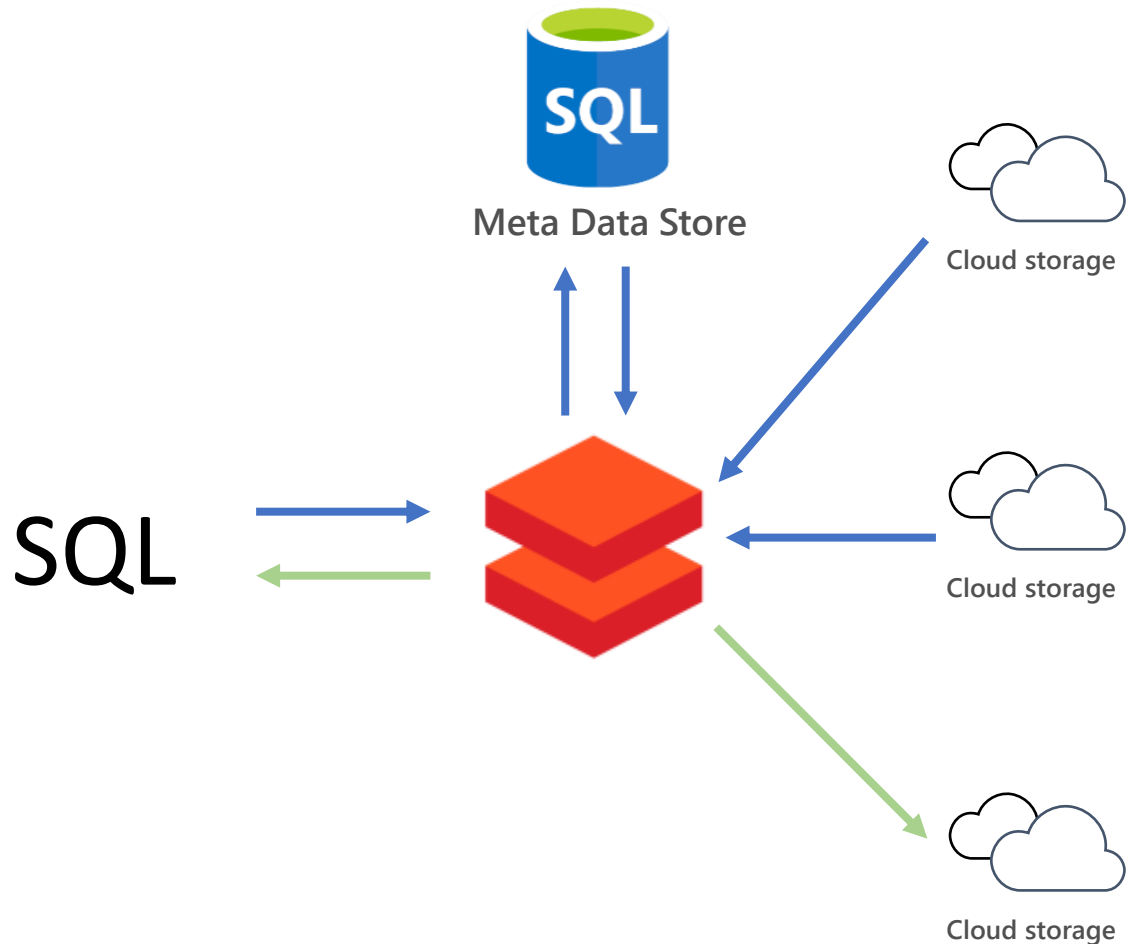
# SQL with Databricks

# When to use Databricks (over RDBMS)

- Scalability
- Flexibility
  - Structured data
  - Unstructured data
- Open Standard
  - Apache Spark
- Single tool for all workloads
  - ETL
  - Data Scientists
  - BI / Reporting
- Cloud solution

Batch processing only – no OLTP!!!

# Processing of a (SQL) query



1. Client submits SQL Query
2. Databricks queries Meta Data Catalog
  - Checks syntax/columns
  - Returns storage locations
3. Databricks queries storage for raw data
4. Data is loaded into memory of nodes
5. Data is processed on nodes using Spark
6. "Result"
  - Data is written directly to storage services
  - Data is returned to client

# The HIVE Meta-Store

- Contains meta-data of all SQL objects
  - Tables (managed vs. external)
  - Views
  - Functions
- Location of the data
- Structure
- Format
- ...

Tables are just references,  
the data resides on the  
storage!

# SELECT and INSERT

## SELECT

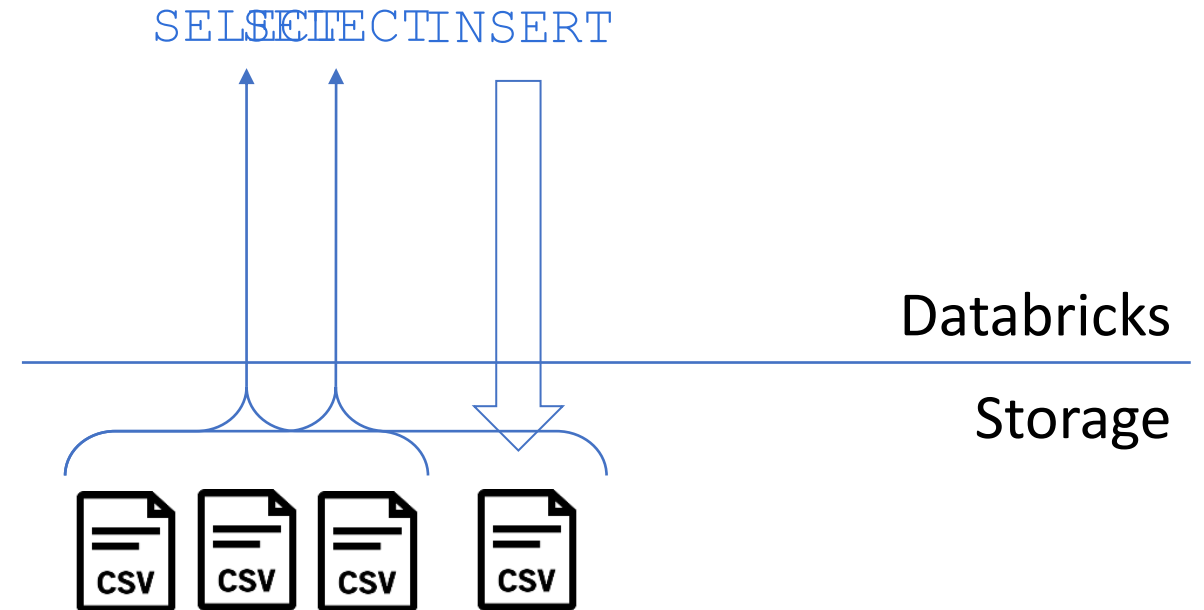
- Reads content of files

## INSERT

- Creates new file on the storage

## UPDATE / DELETE / MERGE

- Not natively support
- → Delta Lake





# Supported SQL features

- ANSI SQL

- Joins
- Groupings/Aggregations
- Rollup/Cube/GroupingSets
- Subselects/CTEs
- Window Functions
- Built-in Functions (date, text, ...)
- ...

- Additional features

- Complex datatypes
  - struct / map / array
- Custom functions
- Custom aggregators
- Caching
- Easy transition to other languages
  - Python, R, Scala, Java, ...

DEMO

Delta Lake

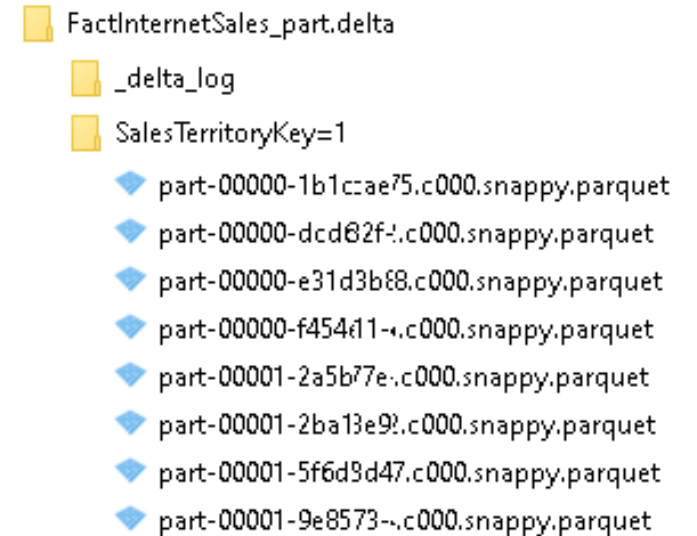
# Delta Lake – [delta.io](https://delta.io)

Delta Lake is an open-source storage layer that brings ACID transactions to Apache Spark™ and big data workloads.

- ACID compliant transactions
  - Optimistic Concurrency Control
- Support for UPDATE / MERGE
- Time-Travel
- Schema enforcement and evolution
  - Across multiple files/folders
- Batch & Streaming
- 100% compatible with Apache Spark































# Delta Lake – [delta.io](https://delta.io)

- Everything is stored in one folder
  - Data
  - Meta-data
  - Transaction log / `_delta_log`
- Could basically Copy & Paste whole Delta table
- Hive meta-store only needs location



# Delta Log

- Contains all meta-data
  - Table schema
  - References to files
- Stored as JSON and Parquet
- Stores all transactions
  - Files added/deleted
  - Changes in meta-data
  - Transaction metric
- Allows concurrency control
- Used for time-travel

Name
 _last_checkpoint
 00000000000000000000.crc
 00000000000000000000.json
 00000000000000000001.crc
 00000000000000000001.json
 00000000000000000002.crc
 00000000000000000002.json
 00000000000000000003.crc
 00000000000000000003.json
 00000000000000000004.crc
 00000000000000000004.json
 00000000000000000005.crc
 00000000000000000005.json
 00000000000000000006.crc
 00000000000000000006.json
 00000000000000000007.crc
 00000000000000000007.json
 00000000000000000008.crc
 00000000000000000008.json
 00000000000000000009.crc
 00000000000000000009.json
 00000000000000000010.checkpoint.parquet
 00000000000000000010.crc
 00000000000000000010.json
 00000000000000000011.crc
 00000000000000000011.json
 00000000000000000012.crc
 00000000000000000012.json
 00000000000000000013.crc
 00000000000000000013.json

# CREATE new Delta Table in Hive meta-store

```
CREATE TABLE IF NOT EXISTS MyTable  
(id INT, name STRING, region INT)  
USING DELTA  
LOCATION '/mnt/adls/tables/DimProductDelta'  
PARTITIONED BY (region)  
TBLPROPERTIES ('myKey' = 'myValue')
```

# Use existing Delta Table in Hive meta-store

```
CREATE TABLE IF NOT EXISTS MyTable  
(id INT, name STRING, region INT)  
USING DELTA  
LOCATION '/mnt/adls/tables/MyTable'  
PARTITIONED BY (region)  
TBLPROPERTIES ('myKey' = 'myValue')
```

Omit everything except

- USING
- LOCATION



# DML Operations – Delta Lake - UPDATE

User

Product	Price
Notebook	900 €
PC	1,500 €
Tablet	500 €

```
UPDATE TABLE DimProduct
SET Price = 1300
WHERE Product = 'PC'
```

Product	Price
Notebook	900 €
PC	1,300 €
Tablet	500 €

\_delta\_log

```
000000000.json
"add": {
  "path": "part-01.parquet",
  ...
}
```

```
000000001.json
"remove": { "path": "part-01.parquet", ... },
"add": { "path": "part-02.parquet", ... }
```

Storage



Parquet  
part-01  
(3 rows)



Parquet  
part-01  
(3 rows)



Parquet  
part-02  
(3 rows)

# DML Operations - Delta Lake – DELETE

User

Product	Price
Notebook	900 €
PC	1,300 €
Tablet	500 €

```
DELETE FROM DimProduct
WHERE Product = 'PC'
```

Product	Price
Notebook	900 €
Tablet	500 €

\_delta\_log

000000000.json

```
"add": {
  "path": "part-01.parquet",
  ... }
```

000000001.json

```
"remove": {
  "path": "part-01.parquet", ... },
"add": {
  "path": "part-02.parquet", ... }
```

000000002.json

```
"remove": { "path": "part-02.parquet", ... },
"add": { "path": "part-03.parquet", ... }
```

Storage



Parquet  
part-01  
(3 rows)



Parquet  
part-02  
(3 rows)



Parquet  
part-01  
(3 rows)



Parquet  
part-02  
(3 rows)



Parquet  
part-03  
(2 rows)

# DML Operations - Delta Lake – INSERT

User

Product	Price
Notebook	900 €
Tablet	500 €

```
INSERT INTO DimProduct
VALUES ('Monitor', 200)
```

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

\_delta\_log

000000000.json

000000001.json

```
"remove": {
  "path": "part-01.parquet", ... },
"add": {
  "path": "part-02.parquet", ... }
```

000000002.json

```
"remove": { "path": "part-02.parquet", ... },
"add": { "path": "part-03.parquet", ... }
```

000000003.json

```
"add": { "path": "part-04.parquet", ... }
```

Storage



Parquet

part-01

(3 rows)



Parquet

part-02

(3 rows)



Parquet

part-03

(2 rows)



Parquet

part-01

(3 rows)



Parquet

part-02

(3 rows)



Parquet

part-03

(2 rows)



Parquet

part-04

(1 row)

# DML Operations - Delta Lake – VACUUM

User

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

VACUUM DimProduct

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

\_delta\_log

000000000.json  
000000001.json

000000002.json  
"remove": {  
  "path": "part-02.parquet", ... },  
"add": {  
  "path": "part-03.parquet", ... }

000000003.json  
"add": {  
  "path": "part-04.parquet",  
  ... }

Storage



Parquet  
part-01  
(3 rows)



Parquet  
part-02  
(3 rows)



Parquet  
part-03  
(2 rows)



Parquet  
part-04  
(1 row)



Parquet  
part-03  
(2 rows)



Parquet  
part-04  
(1 row)

# DML Operations - Delta Lake – OPTIMIZE

User

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

OPTIMIZE DimProduct

Product	Price
Notebook	900 €
Tablet	500 €
Monitor	200 €

\_delta\_log


000000000.json  
000000001.json


000000002.json  
000000002.json

000000003.json

"remove": {  
  "path": "part-03.parquet", "part-04.parquet" },  
"add": {  
  "path": "part-05.parquet", ... }  
}

Storage

  
Parquet  
part-03  
(2 rows)

  
Parquet  
part-04  
(1 row)

  
Parquet  
part-03  
(2 rows)

  
Parquet  
part-04  
(1 row)

  
Parquet  
part-05  
(3 rows)

# UPDATE / DELETE / MERGE

- Always results in new files! Even a DELETE!
- Operations are logged in `_delta_log`
  - Old files are invalidated
  - New files are added/referenced
- Conflicts have to be handled by the User!

Can create A LOT of files!

# OPTIMIZE / VACUUM

## To manage files

### OPTIMIZE

- Collapse small files into bigger files
- Bin-Packing / Ordering
- Improve query performance
- Creates another copy of the data!

```
OPTIMIZE events  
WHERE date = 20200101  
ZORDER BY (eventType)
```

### VACUUM

- Removes unreferenced files older than X days
- Never touches latest version of files

```
VACUUM events  
[RETAIN num HOURS]  
[DRY RUN]
```

DEMO



# Advanced SQL techniques

# Advanced SQL techniques

## User Defined Functions (UDFs)

### UDFs in SQL

- Scalar valued
- Table valued
- Temporary or persisted
- Can be nested

```
CREATE OR REPLACE FUNCTION blue()  
  RETURNS STRING  
  COMMENT 'Blue color code'  
  LANGUAGE SQL  
  RETURN '0000FF'  
;
```

# Advanced SQL techniques

## User Defined Functions (UDFs)

### UDFs from Code

- Python/R/Scala/JAVA
- Scalar valued only (but complex types)
- Must be registered to be used in SQL

### Session-Level only

- Can be nested

### User Defined Aggregates (UDAs)

- SCALA and JAVA only

# Advanced SQL techniques

## External JDBC sources

- Connect to any JDBC source
- Exposed as regular SQL table

```
CREATE TABLE myJdbcTable
USING org.apache.spark.sql.jdbc
OPTIONS (
  url "jdbc:<databaseServerType>://<jdbcHostname>:<jdbcPort>",
  table "<jdbcDatabase>.myTable",
  user "<jdbcUsername>",
  password "<jdbcPassword>"
)
```